


Securing Publish/Subscribe

M.A.Sc. Candidate: Javier Munster

Thesis Supervisor: Hans-Arno Jacobsen


Department of Electrical and Computer Engineering

University of Toronto




Agenda

- Where is pub/sub used?
- What is pub/sub?
- Why use pub/sub?
- Current state-of-the art summary
- HyShare




Pub/Sub in the Wild






Pub/Sub Applications





Security Perspective

- **Confidentiality**
 - Sensitive data is being processed
- **Authorization**
 - Placing limits on who can do what within the system
- **Anonymity**
 - Sensitive data must be analyzed
 - Identity of data sources and sinks must be protected
- **Integrity**
 - Critical infrastructure must always be available
 - Attack resiliency



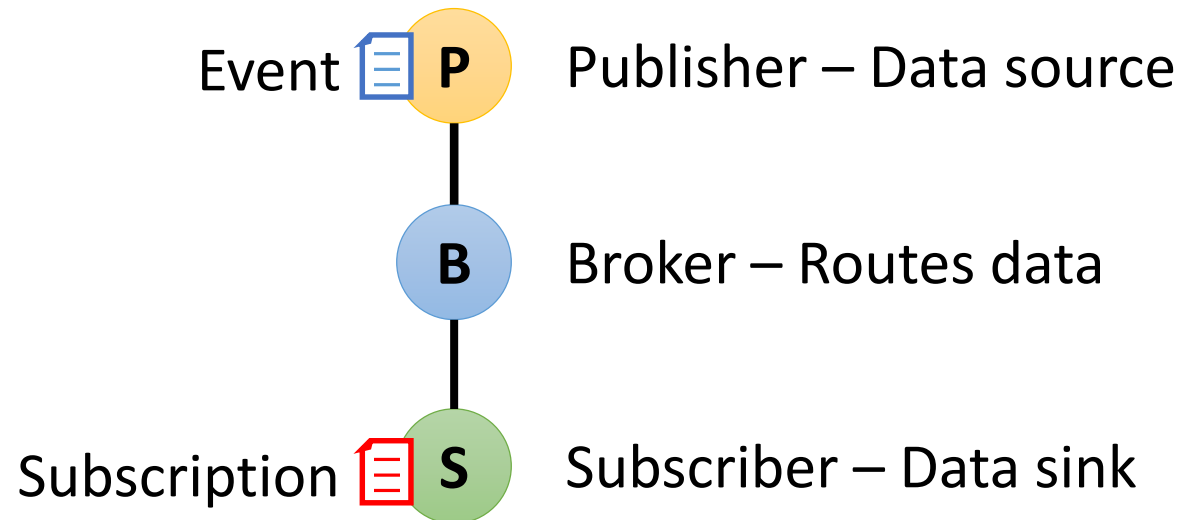
Contributions of Thesis Work

1. Categorization and analysis of the existing pub/sub security research¹
2. Identification and analysis of the limitations of the existing approaches¹
3. HyShare – A novel secret sharing solution used to ensure privacy²


1 – To be published in ACM CSUR

2 – Published in DEBS 2018



What is Publish/Subscribe?




Publish/Subscribe Example



Pub/Sub Terminology






Subscription Anatomy

1. Topic-Based Pub/Sub:

- All events belong to a topic
- Subscriptions specify the topic subscriber is interested in

2. Content-Based Pub/Sub:

- Events include attribute-value pairs
 - Subscriptions contain a subset of attributes subscribers are interested in
- Brokers must match event topic or attributes to subscriptions as part of routing events



Why Use Pub/Sub?

1. Space Decoupling


- Clients do not need to know of one another

2. Time Decoupling

- Sending an event is decoupled from receiving an event


3. Synchronization Decoupling

- Clients do not block when sending or receiving events

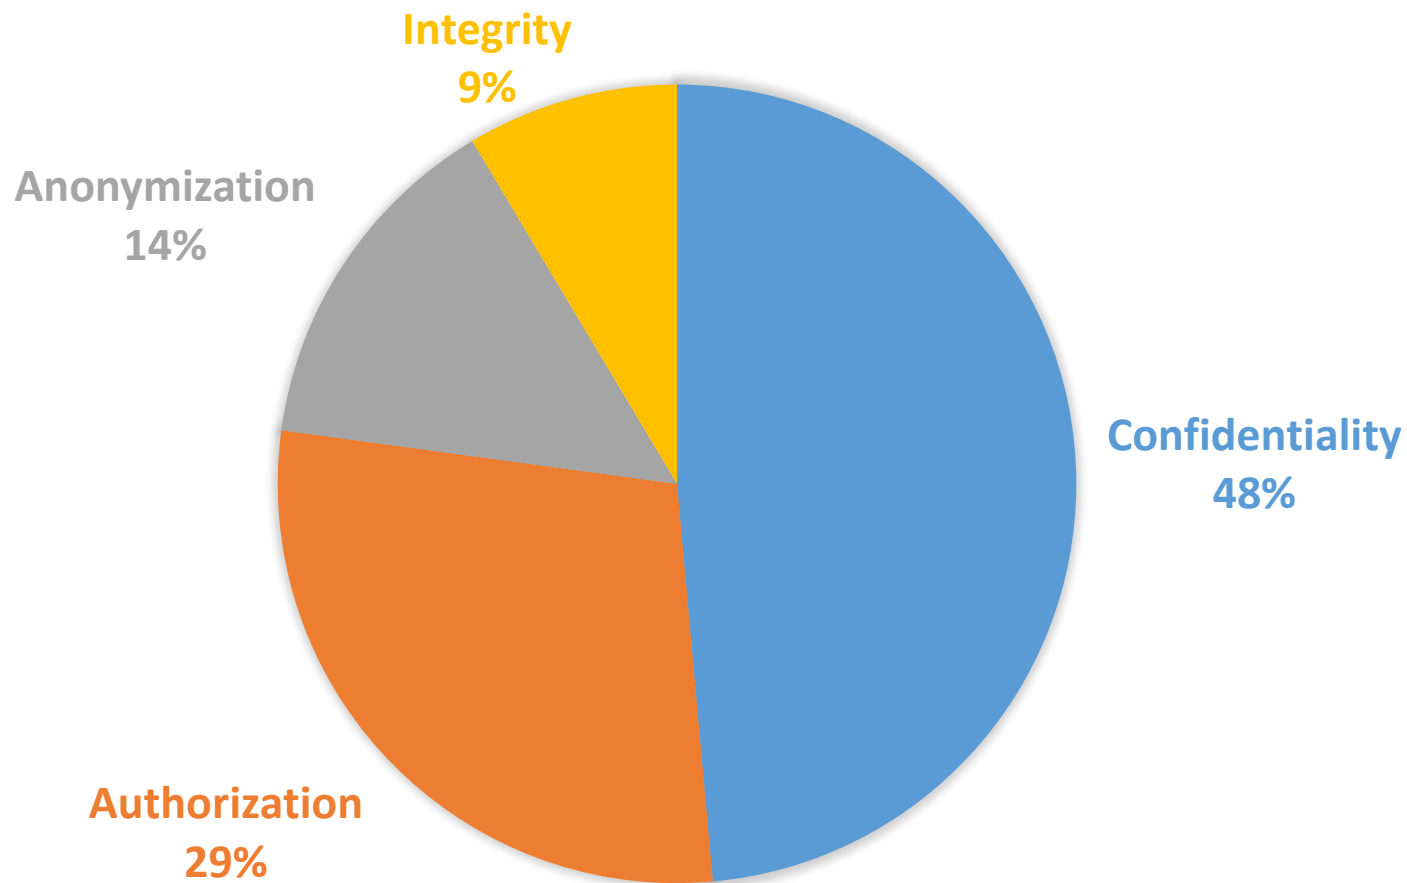



Research Categories

- Confidentiality
 - Hiding the content of events and subscriptions from others
- Authorization
 - Controlling which entities can do what actions within the system
- Anonymization
 - Hiding user identifying information from others
- Integrity
 - Attacks and mitigation strategies for pub/sub




Research Focus





Confidentiality

- Hide sensitive event and subscription information from brokers and external entities
- Brokers still need to be able to filter events based on subscriptions
- Many existing techniques cannot be directly applied in content-based pub/sub
- Techniques:
 - Symmetric-Key Encryption, Homomorphic Cryptosystems, ASPE, Functional Encryption, Multiple Layer Commutative Encryption, Oblivious Transfer, Secret Sharing

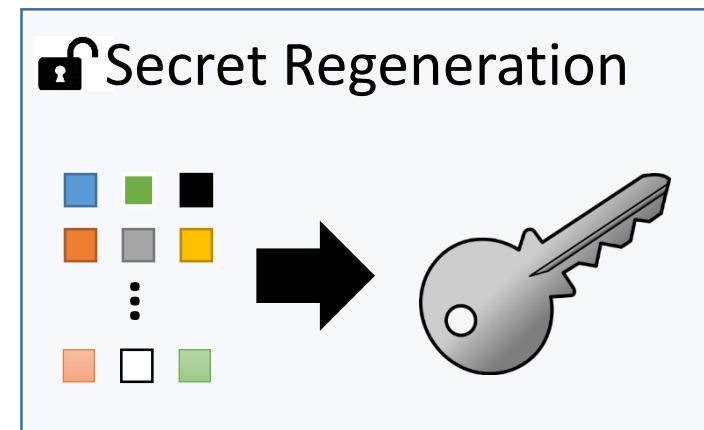



Confidentiality Limitations

- All solutions require the use of a service that is:
 1. Unilaterally trusted
 2. Universally available
 3. Out-of-band
- All solutions use the honest-but-curious broker threat model
- Secret sharing is the only exception but it requires a very large number of messages in order to share a secret

(k, n) -Secret Sharing Scheme


- Based on linear interpolation
- Introduces two operations:



How to Share a Secret, Adi Shamir 1979




Fragment Generation


- Input: Secret s , input parameters k, n
- Output: n fragments
- At least k fragments necessary to regenerate s
- $n \geq k > 1$



Secret Regeneration

- Input: k fragments
- Output: Secret s
- Reverse operation to Fragment Generation






Secret Sharing Propagation Scheme (SSPS)

- Anyone with k fragments can regenerate secret s
- Secret is secure as long as no one ever **sees** more than $k - 1$ fragments
 - No broker ever gets $k - 1$ fragments for a given secret
- In pub/sub, physically redundant delivery paths are used to disseminate fragments
 - Referred to as parallel paths
 - No single broker belongs to k parallel paths
 - ⇔ no single broker can regenerate secret s


SSPS Example 1

- $n = 3$
- $k = 2$ or 3
- Up to $k - 1$ brokers may collude




SSPS Example 2 – Multiple Hops

- Multiple hops between source and sink
- Parallel paths are highlighted
- Parallel paths are of equal length for simplicity




SSPS Example 2 – Collusion Tolerance

- Up to $k - 1$ parallel paths may contain colluding brokers




SSPS Example 2 – Compromised Topology

- Collusion between single broker in each path defeats scheme




Iterative SSPS (ISSPS)

- Brokers apply (k, n) -Secret Sharing Scheme on each received fragment




Iterative SSPS Advantage


- All paths collude without defeating scheme
- Proof published in our DEBS 2018 paper



ISSPS Disadvantage

- Exponential growth in number of messages needed as path length between publisher and subscribers increases






HyShare


- Use hybrid broker networks to share secret
- Leverages recent developments in Intel's SGX
- Hybrid broker network composed of SGX-enabled brokers and ISSPS-brokers
- ISSPS-brokers run ISSPS
- SGX-enabled brokers regenerate the original secret then run ISSPS
- More SGX-enabled brokers means:
 - Fewer messages required to share a secret
 - Less hardware diversity
 - Increased cost

HyShare Simple Example




ISSPS


VS




HyShare




Number of Messages Needed to Share a Secret




Number of Messages Needed to Share a Secret


 $k = 3$

Stock Quote Dissemination

 $k = 3$


Secret Regeneration Overhead






Conclusions

- Categorized the existing pub/sub research
- Identified limitations of existing approaches
- Introduced HyShare, a novel secret sharing scheme that leverages SGX in pub/sub for the first time
- Trusted service used only when bringing up pub/sub network
- No other trusted or out-of-band resources required
- Broker collusion tolerance




Future Work


- Identify what security features are provided by existing pub/sub systems
- Evaluate effects of smarter placement strategies of SGX-enabled brokers
- Examine the effects of leveraging an entire broker network that can use SGX




Questions?




Symmetric-Key Encryption




Homomorphic Cryptosystems




ASPE




Functional Encryption



Multiple Layer Commutative Encryption




Oblivious Transfer




Authorization

- The control of actions being performed within the pub/sub system
- Client Access Control
 - Arbitrary control over messages send or received by client
 - Technique: policy (pre and post matching)
- Broker Access Control
 - Control over messages transmitted within broker network
 - Techniques: Hop-Level Access Control, Domain-Based Access Control




Policy


- Pre- and post-matching of policy rules



Hop-Level Access Control




Domain-Based Access Control




Anonymity


- The hiding of identities and identifying information
- Communication Anonymity
 - The hiding of identities during communication of messages
 - Techniques: Onion Routing, Logical Layer Scheme
- Data Anonymity
 - The sanitization of data being sent of identifying information
 - Techniques: k-anonymity, ℓ -diversity




Onion Routing




Logical Layer Scheme



k-Anonymity




ℓ -Diversity




Integrity


- Attacks and mitigation strategies specifically for pub/sub
- Attacks: Denial of Service, Overlay Scan Attack, Bogus Broker Attack



Denial of Service



Overlay Scan Attack




Bogus Broker Attack


Trusted Execution Environments

- Intel SGX, AMD Platform Security Processor, ARM TrustZone
- Protected area of hardware
- Guarantees data confidentiality and integrity

Normal Attack Surface:



TEE Attack Surface:




Attack Surface



Intel Software Guard Extension (SGX)

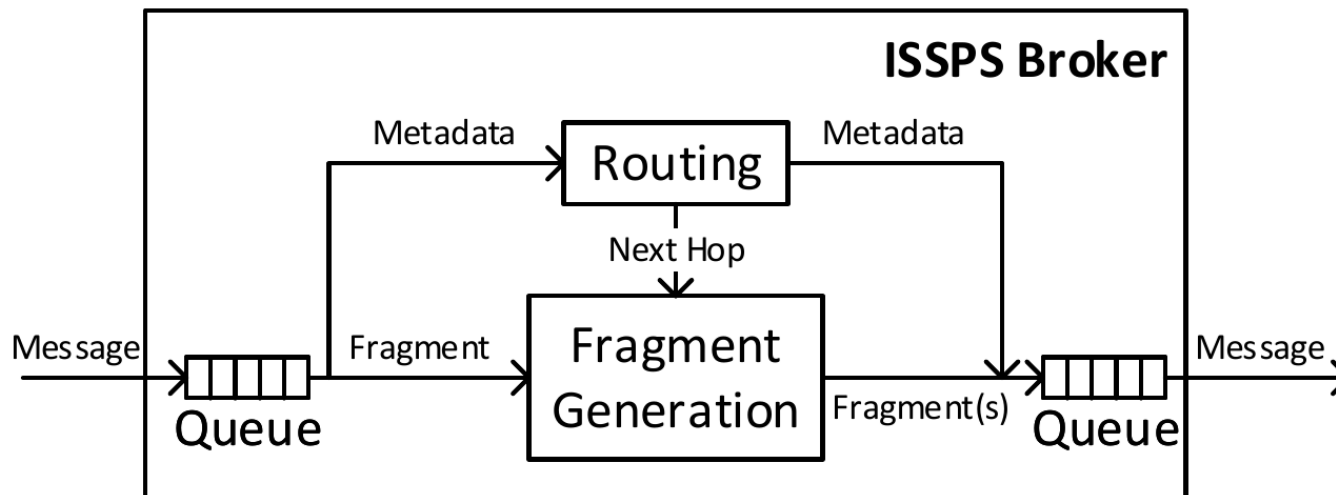
- SGX provides protected containers which:



- Execute code securely
- Provide confidentiality
- Provide integrity
- Have full access to app. data
- Can run attestation protocol

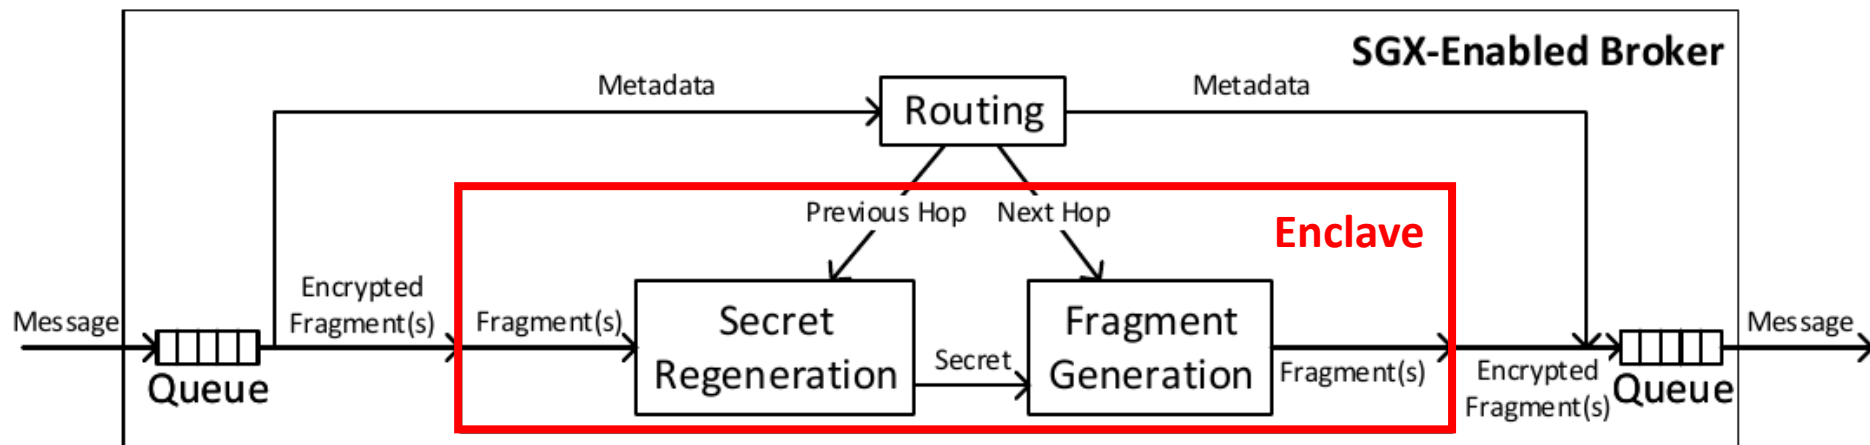
ISSPS-Brokers


- Operate similar to ISSPS
- Generates new fragments if next hop is another ISSPS-broker
- Otherwise, forwards fragments



SGX-Enabled Brokers


- Receives all fragments within enclave
- Regenerates original secret
- Generates n fragments if the next hop is an ISSPS-broker
- Forwards encrypted secret otherwise






Why Share Secrets?

- Keep sensitive information secret from the broker network
 - Personal information, healthcare records, etc.
 - Leak investment strategies
- It's the law
 - Healthcare records must be encrypted before sending
 - General Data Protection Regulation (EU)
- Required by the pub/sub confidentiality literature
 - Homomorphic Cryptosystems, ASPE, Functional Encryption, Multiple Layer Commutative Encryption




Research Problem

- Disseminate a secret/key from a publisher to a subset of subscribers
- Maintain the space decoupling property
- Brokers may share or leak information to learn secret (i.e., collude)
- Minimize trust assumptions required
- Minimize number of messages required



Assumptions of Confidentiality Solutions

- Require the sharing of a key, secret or security parameter(s)
- Rely on a dissemination service that is:
 - Out-of-band
 - Unilaterally trusted
 - Universally available
- Must use dissemination service whenever a subscription event occurs
- Weaken decoupling property between clients



Pure SGX-Enabled Broker Topologies

- Advantages:
 - Conceptually simple to have one type of broker
 - Number of messages needed to share a secret scales linearly with path length
- Disadvantages:
 - Cost of processors that support SGX
 - Lack of SGX adoption in cloud service providers
 - Severe limits on hardware diversity

Parallel Path Generation

